

The Accidental Neuron: How a Forgotten Wire Could Slash AI's Energy Appetite

From: IEEE Spectrum Mined: 2026-07-02 5 min read

Today's AI is a energy hog. Training a single large language model can consume as much electricity as hundreds of homes use in a year, largely because GPUs spend immense power shuttling data between memory and processors. But the human brain runs on about 20 watts less than a dim light bulb while outperforming supercomputers at many tasks. The secret? Neurons and synapses that process information directly where it's stored. For decades, **neuromorphic** engineers have tried to replicate this efficiency by building artificial neurons and synapses. Most efforts either rely on exotic new materials that aren't manufacturable at scale, or they cobble together dozens of standard transistors to emulate a single neuron a bulky, power-hungry approach. Then came a happy accident. In 2024, a student in a lab was testing a memory circuit with a **MOSFET**, a common transistor. He forgot to ground the bulk terminal the fourth connection usually tied to the chip's substrate. What he saw was astonishing: a sudden, nonlinear spike in current when voltage hit a threshold, followed by an automatic relaxation. It was a near-perfect artificial neuron. The phenomenon occurs because a **MOSFET** harbors a hidden **bipolar junction transistor**. When the bulk terminal floats, charge carriers accumulate, activating this **parasitic transistor** and causing a **hysteresis** loop exactly the spike-and-relax behavior of a real neuron. Even better, the same transistor can be tuned to act as a synapse, with adjustable conductance. This means a single, off-the-shelf transistor can serve as either a neuron or a synapse, no complex circuits needed. The implications are huge. Chips built with these "accidental neurons" could be made in standard **CMOS fabs**, scaling up to massive **neuromorphic** systems. They could cut AI power consumption by a factor of a thousand, bringing sophisticated intelligence to low-power devices. It's a classic case of scientific serendipity a forgotten wire that might just rewire the future of computing.

Word Treasure

neuromorphic -- Refers to computer chips designed to mimic the structure and efficiency of the human brain.

MOSFET -- A common type of transistor used as a switch in nearly all modern electronics.

bipolar junction transistor -- An older type of transistor, sometimes hidden inside MOSFETs, that can amplify current.

hysteresis -- A loop-like behavior where a system's output depends on its past states, useful for creating neuron-like spikes.

CMOS fabs -- Factories that manufacture computer chips using a standard process called Complementary Metal-Oxide-Semiconductor.

parasitic transistor -- An unintended transistor that forms inside a device and is usually considered a problem, but here it enables neuron-like behavior.

Background you need

The human brain is incredibly efficient, running on just 20 watts while performing tasks that require warehouse-sized data centers. For years, engineers have tried to copy the brain's design with neuromorphic chips, like IBM's TrueNorth (2014) and Intel's Loihi (2017), but these often required custom materials or complex circuits cobbling together many transistors.

The accidental neuron discovery echoes other famous lab accidents, like Alexander Fleming's discovery of penicillin. MOSFETs (Metal-Oxide-Semiconductor Field-Effect Transistors) are the tiny switches inside every computer chip, with billions made daily in CMOS fabs. A MOSFET's normally hidden parasitic bipolar junction transistor is usually a nuisance, but here it became the key to a simple, elegant artificial neuron.

This means that brain-like chips could be built using the same cheap, mature manufacturing that produces our everyday electronics. It opens the door to neuromorphic systems that scale up massively, potentially bringing sophisticated AI to low-power devices like phones and sensors.

Break it down

LEAD: Contrast between AI's enormous energy appetite and the brain's 20-watt efficiency.

+ - specific: 'The human brain runs on about 20 watts -- less than a dim light bulb.'

BACKGROUND: Decades of neuromorphic efforts either used exotic materials or complex multi-transistor circuits.

ACCIDENTAL DISCOVERY: Student forgot to ground the bulk terminal of a MOSFET.

+ - WHO: A student (unnamed) in a lab in 2024

+ - WHAT: Observed a sudden nonlinear spike and relaxation -- a near-perfect artificial neuron

EXPLANATION: The floating bulk terminal activates a hidden parasitic bipolar junction transistor, creating a hysteresis loop.

DUAL FUNCTION: Same transistor can be tuned to act as a synapse with adjustable conductance.

IMPLICATIONS: Chips can be made in standard CMOS fabs, enabling massive scale and a 1,000x power reduction.

OPEN QUESTION: When will these systems be commercialized, and can they truly rival biological brains?

Why it matters

If AI can run on devices using a thousand times less power, it could enable real-time language translation, health monitoring, and smart assistants without draining batteries or requiring cloud connections, while also reducing the massive carbon footprint of data centers.

Test what you remember

1. What main problem in AI does the accidental neuron discovery address?
 - (A) Slow processing speed
 - (B) High energy consumption
 - (C) Limited memory capacity
 - (D) Lack of accuracy
2. Which terminal of the MOSFET did the student forget to ground?
 - (A) Source
 - (B) Drain
 - (C) Gate
 - (D) Bulk terminal
3. What behavior did the forgotten-wire circuit exhibit?
 - (A) A constant high current
 - (B) A sudden spike in current followed by automatic relaxation
 - (C) A gradual linear increase
 - (D) Complete shutdown
4. What hidden component inside the MOSFET caused this neuron-like spike?
 - (A) A parasitic bipolar junction transistor
 - (B) An extra gate layer
 - (C) A cooling fin
 - (D) A magnetic core
5. By what factor could these accidental neurons cut AI power consumption?
 - (A) Ten times
 - (B) A hundred times
 - (C) A thousand times
 - (D) Ten thousand times
6. What major advantage do these accidental neurons have over previous neuromorphic designs?
 - (A) They use exotic new materials
 - (B) They can be made in standard CMOS fabs
 - (C) They require dozens of transistors per neuron
 - (D) They work only at very low temperatures

Answer key (try the quiz first!): 1: B 2: D 3: B 4: A 5: C 6: B

Questions to think about

1. Positive / Revolutionary: A single standard transistor that acts as both neuron and synapse could slash AI's power consumption by three orders of magnitude. It uses existing manufacturing, so it could scale quickly and bring brain-like computing to small devices worldwide.

2. Negative / Critical: This is still a lab discovery. Proving these transistors work reliably in large systems, with all the variability and noise of a real chip, is a huge challenge. There may be hidden costs in stability or yield that slow commercial adoption.

3. Neutral / Analytical: Serendipity has always played a role in science, and this case mirrors how repurposing a known quirk of MOSFETs can solve a long-standing problem. It highlights that innovation often comes not from brand-new materials but from looking at old technology differently.

4. **Forward-looking:** If engineers can build reliable neuromorphic chips in standard CMOS fabs, the next decade could see AI move from giant data centers into countless low-power edge devices. This might reshape the entire computing industry, making sustainable AI a reality and enabling truly autonomous drones, sensors, and medical implants.

MY THOUGHTS (at least 20 words):